





Aggregations

```
curl localhost:9200/index-name/_search -d '{
  "size": 0,
  "aggs": {
    "most_foos": {
      "terms": {
        "field": "foo.keyword"
      }
    }
  }
}'
```

Term occurrences

terms: by default, most occurrences of a term. Can order by other criteria (including other aggregations)

significant_terms: terms occurring more often in the query results compared to overall. More expensive, may want to use the **sampler** aggregation

Ranges

range: buckets of documents from defined numeric ranges

date_range/ip_range: same as range, but for dates and IPs

histogram/date_histogram: ranges are fixed from an interval

Statistics

```
"aggs": {
  "avg_retweets": {
    "avg": {
      "field": "retweets"
    }
  }
}
```

value_count/min/max/avg/sum of values from a field

percentiles from a numeric field are approximate

cardinality of terms is also approximate

Grouping by nesting aggregations

The following gets the top results, ordered by **_score**, grouped by the value of **bar** (one hit per value).

```
"query": {
  "match": {
    "foo": "f sh"
  }
},
"size": 0,
"aggs": {
  "most_foo": {
    "terms": {
      "field": "bar.keyword",
      "order": {
        "max_score": "desc"
      }
    },
    "aggs": {
      "max_score": {
        "max": {
          "script": {
            "inline": "_score"
          }
        }
      }
    }
  },
  "top_hit": {
    "top_hits": {
      "size": 1
    }
  }
}
```

Document Relationships

Objects

Good for one-to-one relations or when you're searching a single field:

```
curl -XPOST localhost:9200/blog/posts/-d '{
  "title": "Fish & Chips",
  "author": {
    "first_name": "John",
    "last_name": "Smith"
  }
}'
```

Nested

When you need boundaries between objects (e.g. **first_name:jane AND last_name:smith**).

Mapping needs to specify that the parent field is **nested**:

```
"mappings": {
  "posts": {
    "properties": {
      "authors": {
        "type": "nested"
      }
    }
  }
}
```

Documents look like regular objects (even though they're separate Lucene documents):

```
"authors": [
  {
    "first_name": "John",
    "last_name": "Smith"
  },
  {
    "first_name": "Jane",
    "last_name": "Adams"
  }
]
```

Queries (and aggregations) need to be aware of this and do the join:

```
"query": {
  "nested": {
    "path": "authors",
    "query": {
      "match": {
        "authors.first_name": "Jane"
      }
    }
  }
}
```

Parent-child

When updates are frequent and you want to avoid reindexing the whole ensemble (as you would with nested documents). These are completely separate documents, going in different types:

```
"mappings": {
  "authors": {
    "_parent": {
      "type": "posts"
    }
  }
}
```

Children point to parents via the **_parent** field:

```
curl -XPOST localhost:9200/blog/posts/1 -d '{
  "title": "Fish & Chips"
}'
```

```
curl -XPOST localhost:9200/blog/authors?parent=1 -d '{
  "first_name": "John",
  "last_name": "Smith"
}'
```

```
curl -XPOST localhost:9200/blog/authors?parent=1 -d '{
  "first_name": "Jane",
  "last_name": "Adams"
}'
```

Like with nested documents, the query has to specify that a join needs to be done:

```
"query": {
  "has_child": {
    "type": "authors",
    "query": {
      "match": {
        "first_name": "Jane"
      }
    }
  }
}
```